1. **What is a windowed dataset?** 1 / 1 point

- ○ There's no such thing
- ◉ A fixed-size subset of a time series
- ○ A consistent set of subsets of a time series
- ○ The time series aligned to a fixed shape

✓ Correct

2. **What does 'drop_remainder=true' do?** 1 / 1 point

- ○ It ensures that the data is all the same shape
- ○ It ensures that all data is used
- ○ It ensures that all rows in the data window are the same length by adding data
- ◉ It ensures that all rows in the data window are the same length by cropping data

✓ Correct

3. **What's the correct line of code to split an n column window into n-1 columns for features and 1 column for a label**

1 / 1 point

- dataset = dataset.map(lambda window: (window[n-1], window[1]))
- **dataset = dataset.map(lambda window: (window[:-1], window[-1:]))**
- dataset = dataset.map(lambda window: (window[-1:], window[:-1]))
- dataset = dataset.map(lambda window: (window[n], window[1]))

✓ Correct

4. **What does MSE stand for?**

1 / 1 point

- Mean Slight error
- **Mean Squared error**
- Mean Series error
- Mean Second error

5. **What does MAE stand for?** 1 / 1 point

- ○ Mean Average Error
- ○ Mean Advanced Error
- ◉ Mean Absolute Error
- ○ Mean Active Error

✓ Correct

6. **If time values are in time[], series values are in series[] and we want to split the series into training and validation at time 1000, what is the correct code?** 1 / 1 point

- time_train = time[:split_time]
  x_train = series[:split_time]
  time_valid = time[split_time]
  x_valid = series[split_time]

- time_train = time[split_time]
  x_train = series[split_time]
  time_valid = time[split_time]
  x_valid = series[split_time]

- (selected) time_train = time[:split_time]
  x_train = series[:split_time]
  time_valid = time[split_time:]
  x_valid = series[split_time:]

- time_train = time[split_time]
  x_train = series[split_time]
  time_valid = time[split_time:]
  x_valid = series[split_time:]

✓ Correct

7. **If you want to inspect the learned parameters in a layer after training, what's a good technique to use?** 1 / 1 point

- ○ Run the model with unit data and inspect the output for that layer
- ○ Iterate through the layers dataset of the model to find the layer you want
- ○ Decompile the model and inspect the parameter set for that layer
- ◉ Assign a variable to the layer and add it to the model using that variable. Inspect its properties after training

✓ Correct

8. **How do you set the learning rate of the SGD optimizer?** 1 / 1 point

- ○ Use the RateOfLearning property
- ○ You can't set it
- ○ Use the Rate property
- ◉ Use the lr property

✓ Correct

9. **If you want to amend the learning rate of the optimizer on the fly, after each epoch, what do you do?** 1 / 1 point

- ○ **Use a LearningRateScheduler and pass it as a parameter to a callback**
- ○ **Callback to a custom function and change the SGD property**
- ◉ **Use a LearningRateScheduler object in the callbacks namespace and assign that to the callback**
- ○ **You can't set it**

✓ **Correct**